

ML 2023 PROJECT

Using Beamer

Matteo Rofrano, Francesco Lanci Lanci
Group name: Golia

Master degree in Data Science and Business Informatics
emails:
matteo.rofrano@studenti.unipi.it
f.lancilanci@studenti.unipi.it

PROJECT TYPE B

Outline

- 1 Introduction and Objectives
- 2 MONK'S RESULTS
- 3 CUP RESULTS
- 4 Conclusion and appendix

Introduction and Objectives

In the following project, we will test and compare different models, structures and hyperparameters settings in a scientific way in order to find the best model configuration for the task at hand.

The models and algorithms used are:

- Sklearn Neural network with SGD for MONK and CUP
- Keras Neural network with SGD for MONK and CUP
- Pytorch Neural network with SGD for MONK and CUP
- Sklearn Support Vector Machine for MONK and CUP
- Sklearn KNN for MONK and CUP
- XGBoost for MONK and CUP

For the cup we just used the L2 regularization since we noticed that the different NNs tend to not overfit easily. We used MSE for the training and MEE for the cup evaluation.

Method

- Pandas, NumPy, and Sklearn for the preprocessing (transformation and data splitting)
- Sklearn, Keras, PyTorch and xgboost functions for the models
- Sklearn for the model selection (explorative randomized search for the NNs and detailed grid search)
- Matplotlib for the plots

Method

| Implementation | Sklearn | Keras | PyTorch |
|----------------------------|----------------|----------------|----------------|
| Architecture | sequential MLP | sequential MLP | sequential MLP |
| N layers | 1 | 1 | 1 |
| Activation function | tanh/sigmoid | tanh/sigmoid | tanh/sigmoid |
| Activation function output | sigmoid | sigmoid | sigmoid |
| Training algorithm | SGD mini-batch | SGD mini-batch | SGD mini-batch |
| Regularization | L2 only monk3 | L2 only monk3 | L2 only monk3 |
| Initialization | Glorot | Glorot | Glorot |

Table: Structure NNs Monk

| Implementation | Sklearn | Keras | PyTorch |
|----------------------------|----------------|----------------|----------------|
| Architecture | sequential MLP | sequential MLP | sequential MLP |
| N layers | 1/more | 1/more | 1/more |
| Activation functions | tanh/sigmoid | tanh/sigmoid | tanh/ReLU |
| Activation function output | linear | linear | linear |
| Training algorithm | SGD mini-batch | SGD mini-batch | SGD mini-batch |
| Regularization | L2 | L2 | L2 |
| Initialization | Glorot | Glorot | Glorot |

Table: Structure NNs CUP

Method

SVM

- C
- Kernel
- Gamma
- Epsilon

KNN

- N neighbors
- Weights
- Metric

XGBoost

- N estimators
- Max depth
- Learning rate
- Gamma
regularization
- Lambda
regularization

MONK results

| Task | Parameters | MSE (TR/TS) | Accuracy (TR/TS) |
|--------|---|---------------|------------------|
| MONK 1 | $C = 1$, $\epsilon = 0.3$, $\gamma = 0.3$, kernel = polynomial | 0.0805/0.0993 | 100%/100% |
| MONK 2 | $C = 10$, $\epsilon = 0.4$, $\gamma = 0.1$, kernel = rbf | 0.1278/0.1708 | 100%/80.56% |
| MONK 3 | $C = 1$, $\epsilon = 0.3$, $\gamma = 0.4$, kernel = polynomial | 0.0688/0.1102 | 100%/93.52% |

Table: Results SVM

| Task | Parameters | MSE (TR/TS) | Accuracy (TR/TS) |
|--------|--|---------------|------------------|
| MONK 1 | N neighbors = 9, weights = uniform, metric = euclidean | 0.1250/0.1685 | 86.30%/77.31% |
| MONK 2 | N neighbors = 23, weights = distance, metric = euclidean | 0/0.1268 | 100%/78.70% |
| MONK 3 | N neighbors = 5, weights = distance, metric = manhattan | 0/0.0950 | 100%/88.66% |

Table: Results KNN

| Task | Parameters | MSE (TR/TS) | Accuracy (TR/TS) |
|--------|--|---------------|------------------|
| MONK 1 | N estimators = 200, lr = 0.01, max depth = 5, $\gamma = 0.2$, lamda = 0.5 | 0.0133/0.0218 | 100%/100% |
| MONK 1 | N estimators = 300, lr = 0.001, max depth = 10, $\gamma = 0$, lamda = 0.1 | 0.1396/0.1855 | 100%/81.48% |
| MONK 1 | N estimators = 300, lr = 0.3, max depth = 25, $\gamma = 0.2$, lamda = 0.1 | 0.0129/0.0376 | 100%/93.06% |

Table: Results XGBoost

MONK results

| Task | Parameters | MSE (TR/TS) | Accuracy (TR/TS) |
|---------------|---|-------------------|------------------|
| MONK 1 | Units = 4, f = tanh, lr = 0.05, momentum = 0.9, batch size = 20, epochs = 200 | 0.001215/0.00246 | 100%/100% |
| MONK 2 | Units = 4, f = tanh, lr = 0.1, momentum = 0.8, batch size = 20, epochs = 100 | 0.000609/0.000787 | 100%/100% |
| MONK 3 | Units = 4, f = tanh, lr = 0.1, momentum = 0.9, batch size = 20, epochs = 200 | 0.0565/0.0455 | 93.08%/96.76% |
| MONK 3 (reg.) | Units = 4, f = tanh, lr = 0.01, momentum = 0.5, batch size = 20, epochs = 200, l2 = 0.001 | 0.0735/0.0564 | 93.44%/97.22% |

Table: Results Sklearn MLP

| Task | Parameters | MSE (TR/TS) | Accuracy (TR/TS) |
|---------------|---|-----------------|------------------|
| MONK 1 | Units = 5, f = tanh, lr = 0.9, momentum = 0, batch size = 2, epochs = 100 | 0.00014/0.00019 | 100%/100% |
| MONK 2 | Units = 5, f = tanh, lr = 0.1, momentum = 0.8, batch size = 2, epochs = 100 | 0.00020/0.00030 | 100%/100% |
| MONK 3 | Units = 6, f = tanh, lr = 0.9, momentum = 0.8, batch size = 2, epochs = 200 | 0.0643/0.0375 | 93.55%/96.22% |
| MONK 3 (reg.) | Units = 4, f = sigmoid, lr = 0.1, momentum = 0, batch size = 2, epochs = 100, l2 = 0.01 | 0.1042/0.0866 | 93.44%/97.22% |



Table: Results Keras NN

| Task | Parameters | MSE (TR/TS) | Accuracy (TR/TS) |
|---------------|---|-----------------|------------------|
| MONK 1 | Units = 4, f = tanh, lr = 0.3, momentum = 0.5, batch size = 5, epochs = 100 | 0.00088/0.00162 | 100%/100% |
| MONK 2 | Units = 5, f = tanh, lr = 0.3, momentum = 0.6, batch size = 5, epochs = 150 | 0.00020/0.00032 | 100%/100% |
| MONK 3 | Units = 5, f = tanh, lr = 0.1, momentum = 0.7, batch size = 30, epochs = 80 | 0.0466/0.0362 | 94.26%/96.76% |
| MONK 3 (reg.) | Units = 6, f = tanh, lr = 0.01, momentum = 0.7, batch size = 5, epochs = 100, l2 = 0.01 | 0.0680/0.0546 | 93.44%/97.72% |



Table: Results PyTorch NN

MONK results



MONK 1





MONK 2



MONK 3 (no L2)



MONK 3 (L2)



MONK experiments

During the model selection phase we have analyzed the relations between the parameters that leads to overfitting and underfitting. In 1 a comparison between the final model of pytorch with L2 and a model that comes from the randomized search that underfits.

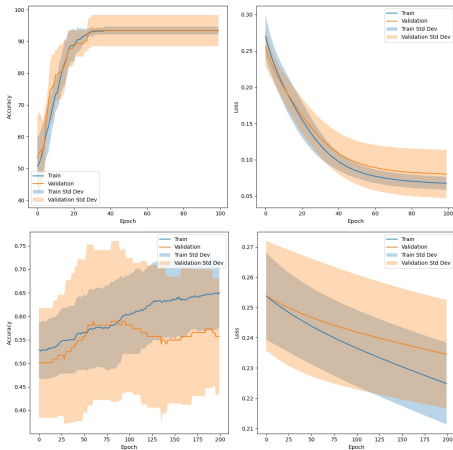


Figure: Loss and accuracy for training and validation with std

CUP validation schema

| Repeated 7-fold-CV | |
|-----------------------|-----|
| Training + Validation | 80% |

| Hold-out | |
|---------------|-----|
| Internal Test | 20% |

Table: Data splitting

- Explorative randomized search for the NNs
- Grid search
- Test of the best model on the internal test set
- Prediction on the blind test set

CUP experiments

The Approach that has lead to the best pytorch NN model is the following:

- starting a randomized search with a high range of parameters' values for a one hidden layer MLP since it is often enough for almost all tasks, and try to find the top 2 best configurations of parameters.
- use the top 2 best configurations to run a new randomized search with higher granularity and return the top 2 models.
- once the parameter for the number of units has converged and we have a small range of values for the other parameters we run a gridsearch.
- repeat this procedure of sequential randomized search + final gridsearch even for a multilayer model.
- compare the one hidden layer model MEE performance with the MEE performance of the multi layer model

CUP results

- PyTorch NN (best model) ranges
- 10 hours for the screening and the grid

| Parameters | Hyperparameters tested |
|----------------------|------------------------|
| N layer | from 1 up to 5 |
| Units for each layer | from 5 up to 250 |
| Activation functions | ReLU, tanh |
| Learning rate | from 0.0001 up to 0.1 |
| L2 | from 0.0001 up to 0.1 |
| Momentum | from 0.1 up to 0.9 |
| Batch size | from 5 up to 100 |
| Epochs | from 50 up to 300 |

Table: Screening phase

CUP results

At the end the best model to use is the one with the lowest MEE on the validation data. Then it has been tested on the internal test data.

| Parameters | Final Grid-search best model |
|------------------------|------------------------------|
| N layer | 2 layers |
| Units first layer | 180, 200 |
| Units for second layer | 20, 40 |
| Activation functions | ReLU |
| Learning rate | 0.001, 0.005, 0.0005 |
| Momentum | 0.9 |
| L2 | 0.0001 |
| Batch size | 20, 40 |
| Epochs | 50, 100 |

Table: Values of parameters of best model

BEST PARAMETERS= 2 layers, 200 units first layer, 40 units second layer, momentum 0.9, learning rate 0.001, L2 0.0001, batch size 20, epochs 100

CUP results

| MEE TR/VL | MEE TR/TS |
|----------------------------------|-------------|
| 0.813 std=0.041/1.0360 std=0.168 | 0.727/0.952 |

Table: Best model results





Figure:
MSE and
MEE for
training and
internal test

Conclusions

During the development of the project, we were able to put into practice the knowledge we acquired during the course, and we became confident in using several libraries. We became more practiced in model selection and assessment, more specifically, we saw how important a screening phase (randomized search in our case) is before the grid search, as it is too time-consuming if used with too many parameter combinations. In addition to a purely computational issue, the screening phase was also useful for us to understand what combinations of parameters did not get along and the direct effects on the loss and accuracy graphs.

Conclusions

Nickname: Golia

Bind test results: Golia_ML-CUP23-TS.csv

We agree to the disclosure and publication of our names, and of the results with preliminary and final ranking.

References

<https://pandas.pydata.org/>

<https://numpy.org/>

<https://scikit-learn.org/stable/>

<https://matplotlib.org/>

<https://keras.io/>

<https://pytorch.org/>

<https://xgboost.readthedocs.io/en/stable/>

Appendix




Figure: Low lr: MONK 2, batch = 2, lr = 0.001, momentum = 0.6

Appendix




Figure: MONK 1: training and validation curves of the best model (Keras)

Appendix




Figure: MONK 2: training and validation curves of the best model (Keras)

Appendix




Figure: CUP: MSE and MEE with std on training and validation